

# Stochastic Processes and Long-Run Behavior

This investigation studies long-run behavior in stock series using volatility clustering, ARCH(2) simulation, and regime-based forecasting.

```
set.seed(1234)
data("EuStockMarkets")
```

## Volatility clustering and ARCH(2) simulation

We analyze each index by: 1) computing log returns and drift, 2) checking ACF for returns and squared returns, 3) fitting an ARCH(2) proxy using AR(2) on squared returns, and 4) simulating a synthetic index with the same length and drift.

```
analyze_arch2 <- function(index_series, label) {
  log_returns <- diff(log(index_series))
  drift <- mean(log_returns)

  par(mfrow = c(1, 2))
  acf(log_returns, main = paste(label, "log returns ACF"))
  acf(log_returns^2, main = paste(label, "squared returns ACF"))

  fit <- arima(log_returns^2, order = c(2, 0, 0))
  phi <- fit$coef[1:2]
  a0 <- fit$coef[3]

  n <- length(index_series)
  y <- numeric(n)
  y[1:2] <- log_returns[1:2]
  z <- rnorm(n)

  for (i in 3:n) {
    s <- sqrt(a0 + phi[1] * y[i - 1]^2 + phi[2] * y[i - 2]^2)
    y[i] <- s * z[i]
  }

  y <- y + drift
  y <- c(log(index_series[1]), y)
  sim_index <- exp(cumsum(y))

  par(mfrow = c(1, 2))
  ts.plot(index_series, main = paste(label, "Index"), ylab = "Index")
  ts.plot(sim_index, main = paste(label, "Simulated Index"), ylab = "Index")

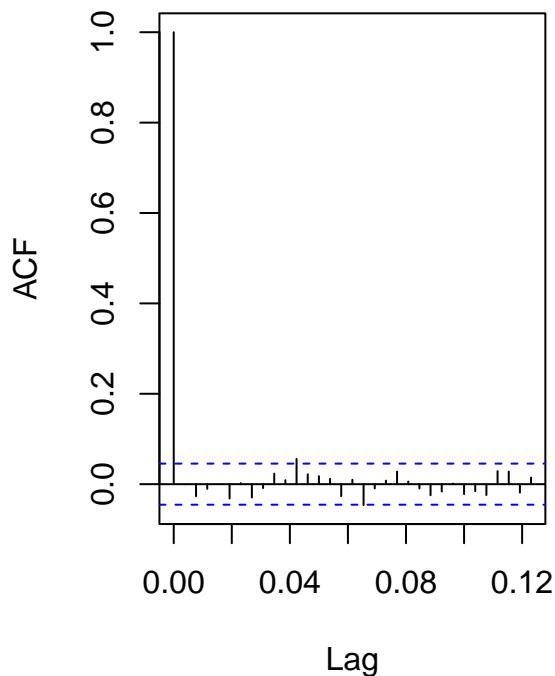
  invisible(list(
    log_returns = log_returns,
    drift = drift,
    fit = fit,
    sim_index = sim_index
```

```
)  
}
```

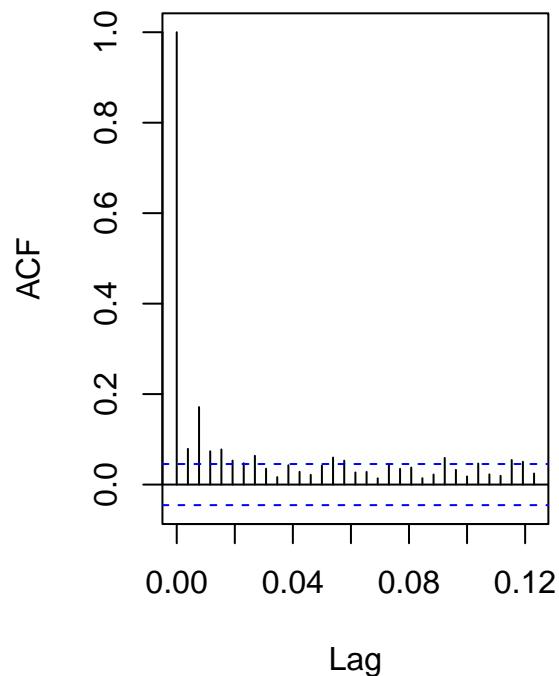
DAX

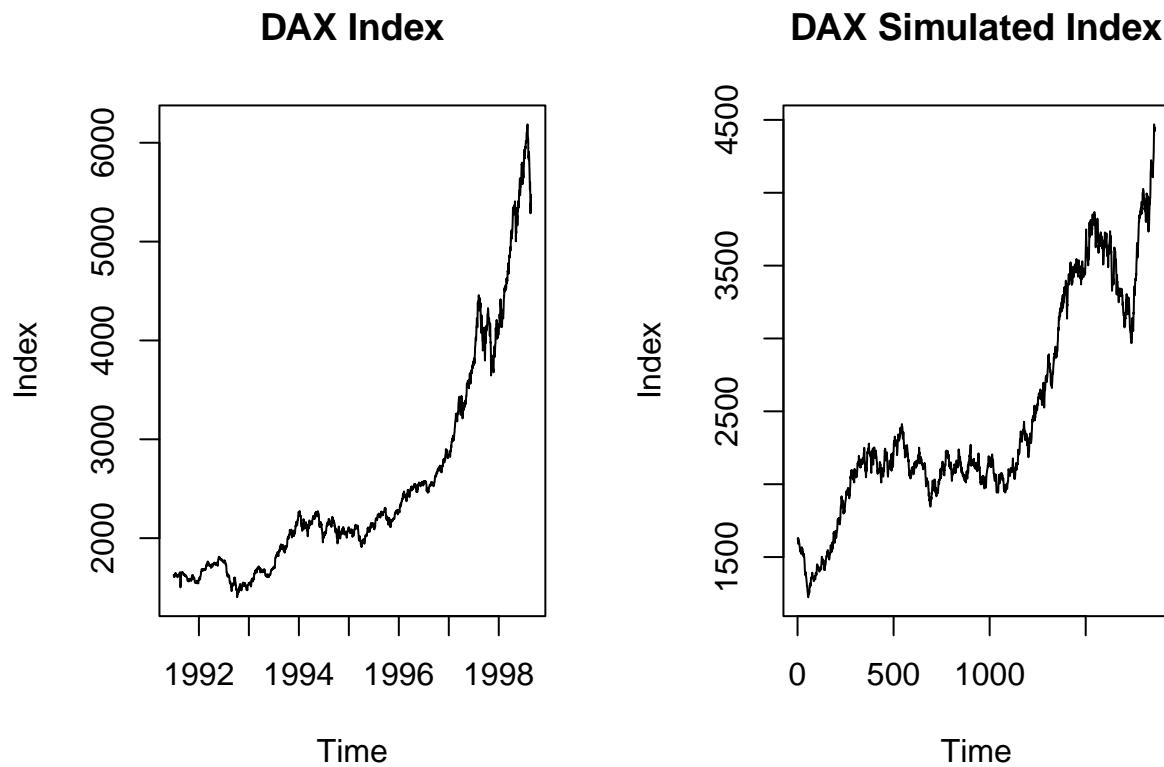
```
dax <- EuStockMarkets[, "DAX"]  
dax_out <- analyze_arch2(dax, "DAX")
```

**DAX log returns ACF**



**DAX squared returns ACF**



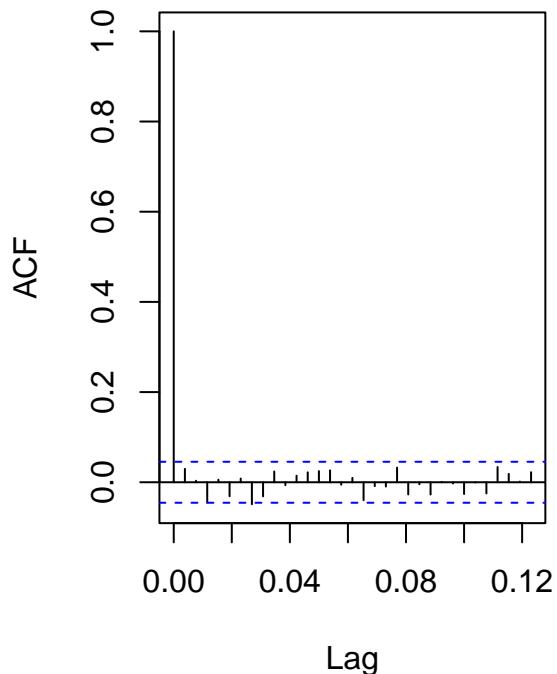


The DAX log-return ACF does not show meaningful spikes beyond lag 0, so there is no evidence of linear dependence. The squared returns ACF shows significant early-lag autocorrelation, indicating volatility clustering.

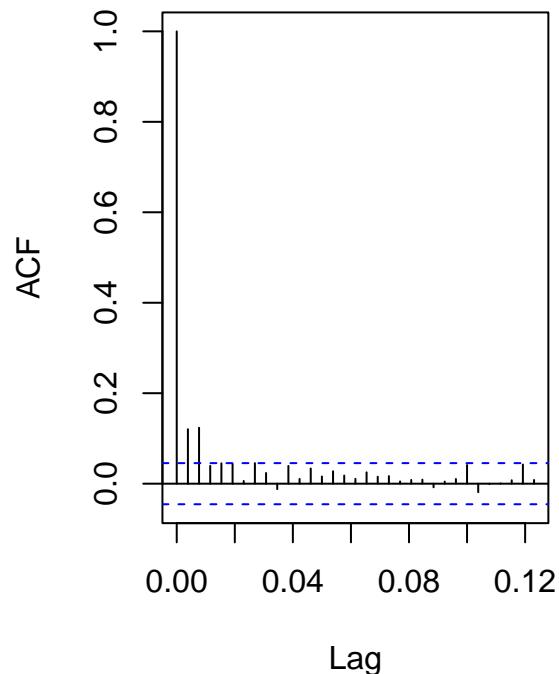
### CAC

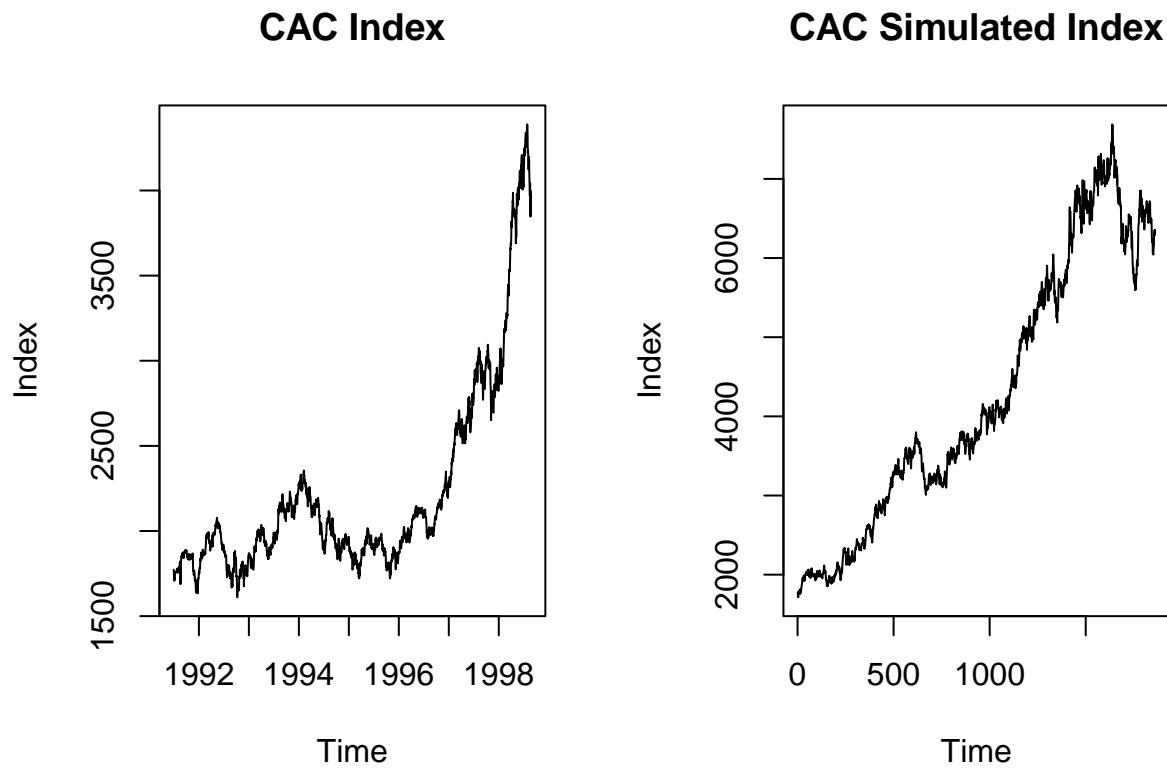
```
cac <- EuStockMarkets[, "CAC"]
cac_out <- analyze_arch2(cac, "CAC")
```

**CAC log returns ACF**



**CAC squared returns ACF**



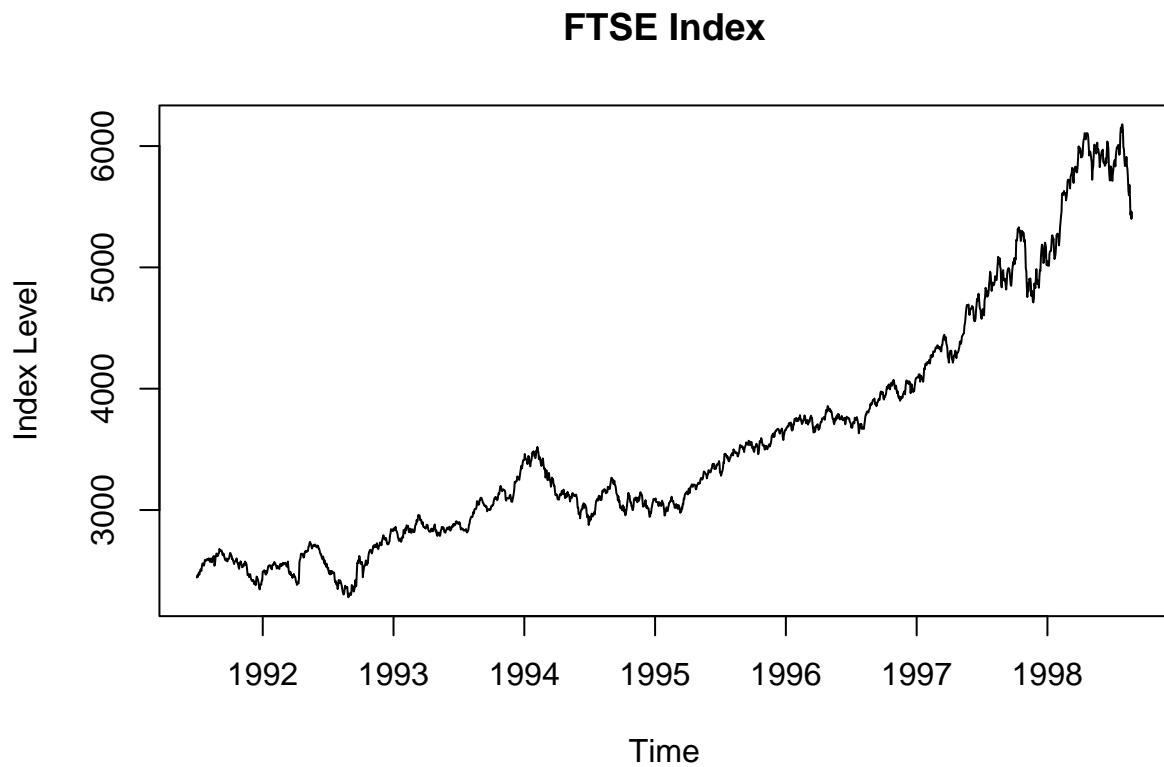


### Regime-based modeling and long-run forecasting (FTSE)

We discretize FTSE log returns into 10 quantile bins, fit a 4-state HMM, and simulate the remaining path to estimate the final index value.

```
ftse <- EuStockMarkets[, "FTSE"]
n <- length(ftse)
log_returns <- diff(log(ftse))
log_returns_q <- c(
  min(log_returns) * 1.01,
  quantile(log_returns, (1:9) / 10),
  max(log_returns) * 1.01
)
log_returns_cut <- cut(log_returns, log_returns_q)
levels(log_returns_cut) <- 1:10

ts.plot(ftse, main = "FTSE Index", ylab = "Index Level", xlab = "Time")
```



```

ftse[995:1005]
## [1] 3199.9 3209.3 3214.9 3226.2 3217.6 3216.7 3220.4 3248.2 3262.6 3264.3
## [11] 3251.7
ftse[n]
## [1] 5455
P <- matrix(c(1:4, 4:1, 1, 4, 4, 1, 4, 1, 4, 1), nrow = 4, byrow = TRUE) / 10
PE <- matrix(
  c(1:10, 10:1, rep(5.5, 10), c(rep(1, 5), rep(10, 5))),
  nrow = 4,
  byrow = TRUE
) / 55

library(HMM)
states <- 1:4
symbols <- 1:10
hmm0 <- initHMM(
  States = states,
  Symbols = symbols,
  startProbs = rep(1 / 4, 4),
  transProbs = P,
  emissionProbs = PE
)
n_training <- 1000

```

```

hmm_fit <- baumWelch(hmm0, observation = log_returns_cut[1:n_training])

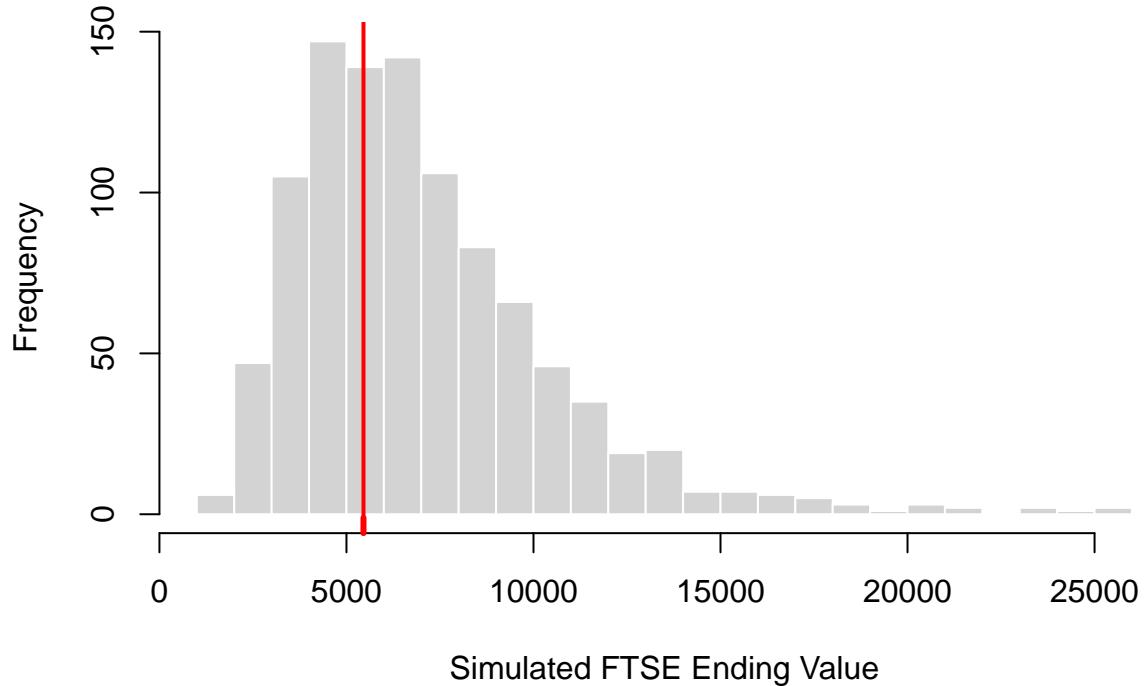
N <- 1000
n_testing <- n - n_training
index_859 <- numeric(N)

for (i in 1:N) {
  index <- as.numeric(simHMM(hmm_fit$hmm, n_testing)$observation)
  index_859[i] <- exp(
    cumsum(
      c(
        log(ftse[n_training]),
        runif(n_testing, min = log_returns_q[index], max = log_returns_q[index + 1])
      )
    )
  )[n_testing]
}

hist(
  index_859,
  breaks = 30,
  main = "Simulated Estimates of FTSE[1859]",
  xlab = "Simulated FTSE Ending Value",
  col = "lightgray",
  border = "white"
)
actual_val <- ftse[n]
rug(actual_val, col = "red", lwd = 3)
abline(v = actual_val, col = "red", lwd = 2)

```

## Simulated Estimates of FTSE[1859]



```
actual_val
```

```
## [1] 5455
```

The simulated values are realistic: the true final value lies within the dominant region of simulated outcomes, even with a heavy upper tail.

### Option pricing (30-day horizon)

We price a European call option with strike 5000 when purchased 30 days before expiry.

```
n_training <- 1829
n_testing <- 1859 - 1829
index_option <- numeric(N)

for (i in 1:N) {
  index <- as.numeric(simHMM(hmm_fit$hmm, n_testing)$observation)
  sim_log_ret <- runif(
    n_testing,
    min = log_returns_q[index],
    max = log_returns_q[index + 1]
  )
  sim_path <- exp(cumsum(c(log(ftse[n_training]), sim_log_ret)))
  index_option[i] <- sim_path[n_testing + 1]
}

K <- 5000
option_price <- mean((index_option - K) * (index_option > K))
```

```
option_price
```

```
## [1] 1142.766
```

With only 30 days to maturity, the expected payoff is much lower than longer-horizon pricing because there is less time for large upward moves.